



AUTONOMOUS EXPLORATION AGENT

Authors:

Federico Bottoni
Nassim Habbash

Task

Definition of a
single-agent
obstacle-ridden,
procedurally generated
environment

Training of the agent
with Proximal Policy
Optimization for a
target localization
objective

Evaluation and
comparison of the
agent's **model**
variations


Tools

Environment and agent
architecture



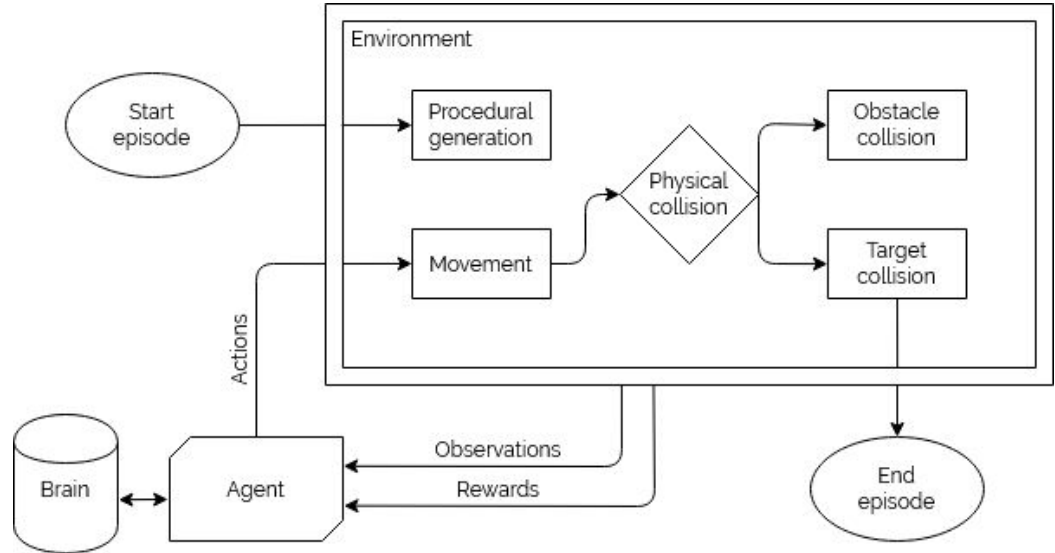
Reinforcement Learning

Unity's ML-Agents

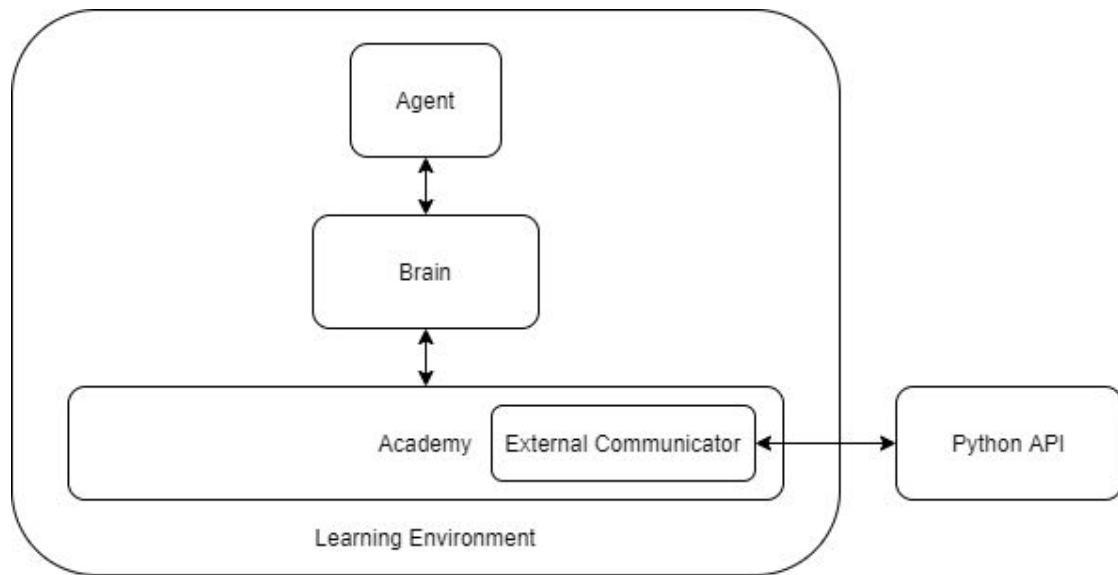


System Architecture

Schema illustrating the **interactions** between the main **actors** of the system



System Architecture



ML-Agent Learning
pipeline


Simulation Main Phases

- ▶ **Environment procedural generation**
 - Parametrized spawning and positioning of the **Agent**, **Obstacles** and **Target**
- ▶ **Agent's Actions and Movement**
 - Agent movement controller built with Unity's **Physics System**
 - Actions inferred by ML-Agents **Brain** interface
 - Action space: **3D Discrete** (2D maneuvering case)
- ▶ **Collisions**
 - **Penalty** function for obstacle collision
 - **Reward** function and episode completion for target collision

Environment Procedural Generation

Constrained generation of the scene ensuring an approximately uniform distribution of the object in the environment according to different parameters, such as:

- Number of obstacles
- Minimum Target distance from the Agent
- Minimum distance between objects in the environment



Environment Procedural Generation

Agent parameters

Dimensions	1x1x1
Max linear velocity	5
Max angular velocity	$5/3\pi$

Environment area parameters

Level area	50x50
Obstacle dimensions	8x8x8
Target dimensions	8x8x8

Sensor parameters

# LIDAR	14
Maximum range	20
Field of view	$[-2/3\pi, 2/3\pi]$

3D Sensor parameters

# LIDAR	42
Maximum Range	40
Horizontal field of view	$[-2/3\pi, 2/3\pi]$
Vertical field of view	$[-1/3\pi, 1/3\pi]$

Static environmental parameters

2D Maneuvering System



Movement

- Simple vehicle physics **approximation**
- Movement through physically grounded **force application**
- **Instant velocity** change to the agent's body (ForceMode.VelocityChange)
- Soft clamping to the max velocity

Actions

- **Decisions** coming from the Brain (model)
- **3D Action Space:**
 - x, z axis **translation** movement
 - yaw axis **rotation** movement

3D Maneuvering System



Movement - Physics system:

- **No gravity**

Actions - Augmented to a **5D** action space:

- x, y, z axis **translation**
- yaw and pitch **rotation**

Learning System

- Uses **Proximal Policy Optimization** as a RL algorithm
- The **reward signal** defines the goal of the task
- **Curriculum learning** scales the difficulty of the task according to the **cumulative reward** reached by the agent
- The same agent's **Brain** is trained on **parallel environments**

Reward Signal - Intrinsic Reward


Given in response to the actions made by the agent in the environment, comprised of two components:

- Reward: *Positive reward - Penalty function*
 - Penalty function: $\alpha * \text{obstacle_collisions} + \beta * \text{time}$
 - Positive reward: $\gamma * \text{target_collisions}$

Reward Signal - Extrinsic Reward

Representing the **curiosity** of the agent.

The more **unexpected** the action taken by the agent, the **higher** the curiosity signal.




Curriculum Learning

Dynamically change the parameters **during training**, making the task **progressively** harder

In some cases, allows for **faster policy convergence**

Curriculum parameters
Reward threshold
Number of obstacles
Minimum spawn distance
Target-Agent distance
Penalty offset

Parallel Learning



Increases the experience throughput of the Agent through parallel instances sharing the Brain

Experiments

Performance analysis of different scenarios:

- Baseline
- Curriculum learning
- Harder Penalty function
- Camera sensors
- 3D maneuvering environment

Experiments - Evaluation

Evaluation made through two different types of measurements:

Traditional RL performance metrics

- Cumulative reward
- Policy loss
- etc

Environmental performance metrics

- Collisions per minute (CPM)
- Targets per minute (TPM)
- Collisions per target (CPT)

Experiments - Baseline

Fixed parameters:

Number of obstacles	10
Min spawn distance	2
Target distance	45

Penalty function:

$$p = \text{collisions} * 0.1 + \text{time} * 0.001$$

Observation sensors: LIDAR

2D maneuvering environment

Experiments - Curriculum

⇒ Curriculum parameters:

Reward thresholds	1	2	2.5	2.8	3	3.5	4
Number of obstacles	8	10	13	15	17	18	20
Min spawn distance	6	6	4	4	3	3	2
Target distance	25	28	30	33	35	37	40

Penalty function:

$$p = \text{collisions} * 0.1 + \text{time} * 0.001$$

Observation sensors: LIDAR

2D maneuvering environment

Experiments - Harder Penalty

⇒ Curriculum parameters:


Reward thresholds	1	2	2.5	2.8	3	3.5	4
Number of obstacles	8	10	13	15	17	18	20
Min spawn distance	6	6	4	4	3	3	2
Target distance	25	28	30	33	35	37	40
Penalty offset	0.5	1.5	2	2.5	2.5	2.5	2.5

⇒ Penalty function:

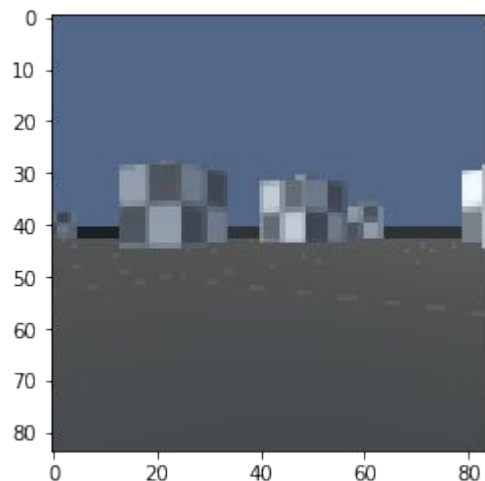
$$p = p_{offset} + collisions + time * 0.001$$

Observation sensors: LIDAR

2D maneuvering environment



Experiments - Camera Sensors



⇒ Curriculum parameters:

Reward thresholds	1	2	2.5	2.8	3	3.5	4
Number of obstacles	8	10	13	15	17	18	20
Min spawn distance	6	6	4	4	3	3	2
Target distance	25	28	30	33	35	37	40

Penalty function:

$$p = \text{collisions} * 0.1 + \text{time} * 0.001$$

⇒ Observation sensors: Camera, 84x84 RGB 2D maneuvering environment

Experiments - 3D maneuvering env.

⇒ Curriculum parameters:

Reward thresholds	1	2	2.5	2.8	3	3.5	4
Number of obstacles	8	10	13	15	17	18	20
Min spawn distance	6	6	4	4	3	3	2
Target distance	25	28	30	33	35	37	40

Penalty function:


$$p = collisions * 0.1 + time * 0.001$$

⇒ Observation sensors: (Extended) LIDAR set

⇒ 3D maneuvering environment


Experiments - Structured environment transferability

Performance evaluation of the **best** between the aforementioned models in a **structured** environment, performing the same task




Results - Performance measures

Environment/Cumulative Reward




Results - Performance measures




Results - Performance measures

Losses/Policy Loss




Results - Performance measures

Policy/Extrinsic Reward



Policy/Curiosity Reward



Results - Performance measures

Training time is not equal between experiments due to training time constraints.

The performance measures show **different levels of convergence** of the models and **different advancements in the lessons** of the curriculum that seem to not correlate directly with the training time.

Direct evaluation on these measures is **hard**:

- Which is the best model? The one with a better reward convergence? The one that advanced to the last lesson? The one with a better policy loss convergence? A mix of all them?
- Which measures are significative of the model performance in the experimental scenario?

Results - Environmental performance

Environmental performance measures allow us to measure the **empirical performance** of the models for the particular task they've been trained for.


Environmental performance is hence a **domain-coupled measure**.

The measures used for this project are:

- **TPM:** Targets per minute
- **CPM:** (Obstacle) Collisions per minute
- **CPT:** Collision per target (CPM/TPM)

Each model has been evaluated on the **same environment**, having the **same parameters**, to make the comparison fair.

Results - Environmental performance



Best overall: *curriculum learning model*.

Almost every model managed to stay **below 2 CPT**.

The *harder penalty model* **did not improve** the performance.

The *camera model* **did not perform** as well as the LIDAR model, but also **did not train as much**.

The *3D maneuvering model* reached almost 2 TPM and CPM, but did not complete the whole curriculum

Conclusions

For robotic locomotion and target localization tasks **reinforcement learning** and **curriculum learning** perform effectively.

The results show a different outcome based on the **evaluation method** chosen, underlining the difficulty to evaluate correctly reinforcement learning scenarios.

The models proposed seems to converge on a policy of **random search**, a behaviour shared between every experiment model. Most probably due to the **lack of memory** of the NN and **procedural generation** of the environment.

The models are able to generalize **obstacle avoidance**.

Conclusions - Future works

- **Add memory to the system:** adding an RNN module (akin to the curiosity module) might allow the model to form an experience buffer of sorts and adopt smarter search policies.
- **Rework the reward and penalty functions:** implement more complex/effective functions, (eg soft-collisions)
- **Compare different RL algorithms**
- **Extend the types of scenarios**

Thanks for listening!

Hello!

I am Jayden Smith

I am here because I love to give presentations.

You can find me at:

@username



1.

TRANSITION HEADLINE

Let's start with the first set of slides



Quotations are commonly printed as a means of inspiration and to invoke philosophical thoughts from the reader.

This is a slide title

- ▷ Here you have a list of items
- ▷ And some text
- ▷ But remember not to overload your slides with content

Your audience will listen to you or read the content, but won't do both.



Big concept

Bring the attention of your audience over a key concept using icons or illustrations

In two or three columns

Yellow

Is the color of gold, butter and ripe lemons. In the spectrum of visible light, yellow is found between green and orange.

Blue

Is the colour of the clear sky and the deep sea. It is located between violet and green on the optical spectrum.


Red

Is the color of blood, and because of this it has historically been associated with sacrifice, danger and courage.



Want big impact?
Use big image.

Use diagrams to explain your ideas



And tables to compare data

	A	B	C
Yellow	10	20	7
Blue	30	15	10
Orange	5	24	16

Maps





89,526,124\$

That's a lot of money



185,244 users

And a lot of users



100%

Total success!

Our process is easy

First

Lorem ipsum dolor sit
amet, consectetur
adipiscing elit.

Second

Lorem ipsum dolor sit
amet, consectetur
adipiscing elit.

Last

Lorem ipsum dolor sit
amet, consectetur
adipiscing elit.

Let's review some concepts



Yellow

Is the color of gold, butter and ripe lemons. In the spectrum of visible light, yellow is found between green and orange.



Yellow

Is the color of gold, butter and ripe lemons. In the spectrum of visible light, yellow is found between green and orange.



Blue

Is the colour of the clear sky and the deep sea. It is located between violet and green on the optical spectrum.



Blue

Is the colour of the clear sky and the deep sea. It is located between violet and green on the optical spectrum.



Red


Is the color of blood, and because of this it has historically been associated with sacrifice, danger and courage.



Red

Is the color of blood, and because of this it has historically been associated with sacrifice, danger and courage.

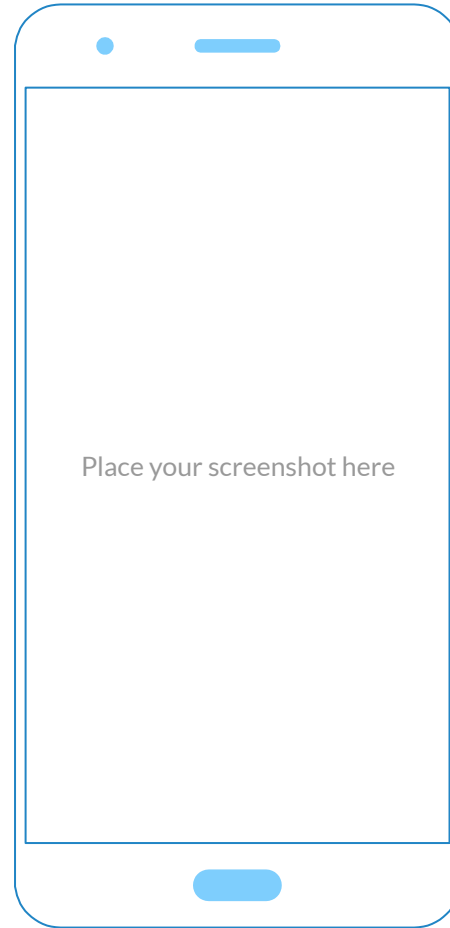
■ A ■ B ■ C



You can copy&paste graphs from [Google Sheets](#)

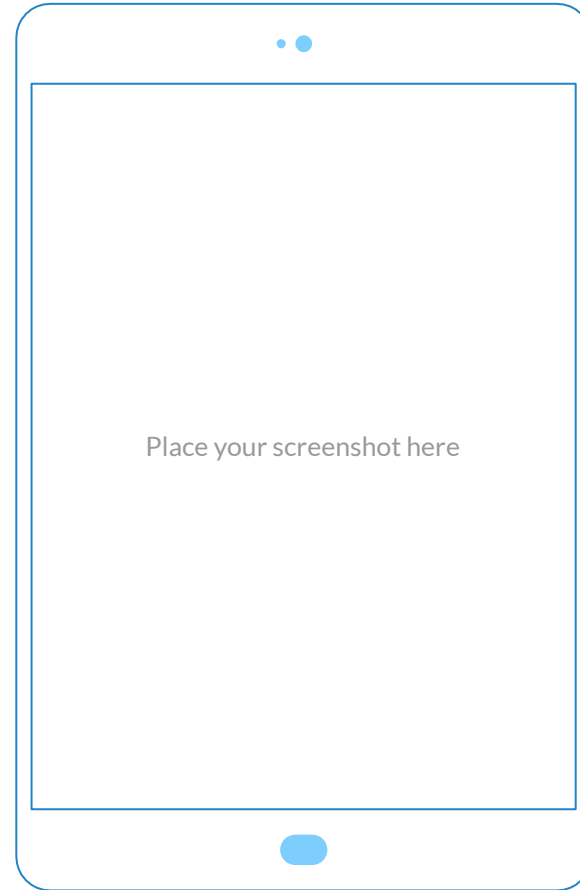
Mobile project

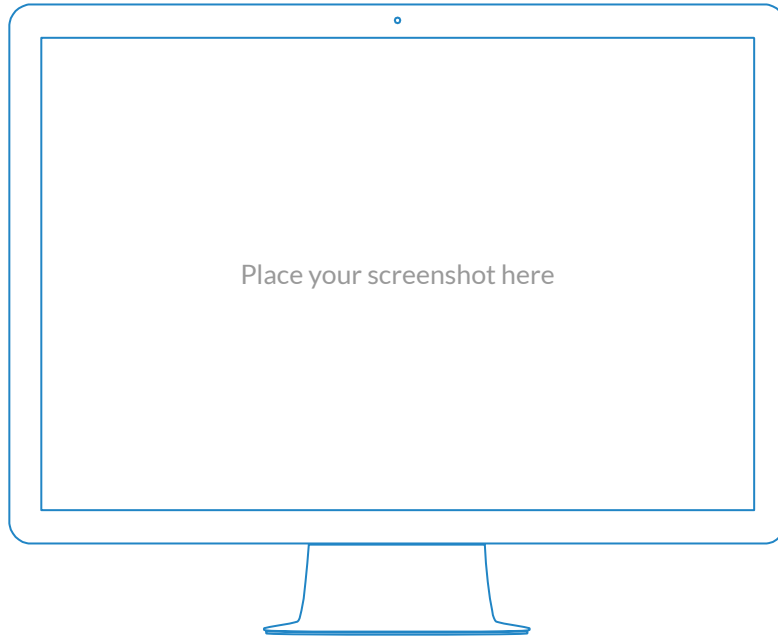
Show and explain your web, app or software projects using these gadget templates.



Tablet project

Show and explain your web, app or software projects using these gadget templates.





Desktop project

Show and explain your web, app or software projects using these gadget templates.

Thanks!

Any questions?

You can find me at:

@username

user@mail.me

Credits

Special thanks to all the people who made and released these awesome resources for free:

- ▷ Presentation template by SlidesCarnival
- ▷ Photographs by Unsplash

Presentation design

This presentations uses the following typographies and colors:

- ▷ Titles: **Raleway**
- ▷ Body copy: **Lato**


You can download the fonts on these pages:

<https://www.fontsquirrel.com/fonts/raleway>

<https://www.fontsquirrel.com/fonts/lato>

- ▷ Dark blue **#2185c5**
- ▷ Light blue **#7ecef**
- ▷ Yellow **#ff9715**
- ▷ Magenta **#f20253**
- ▷ Dark gray **#677480**
- ▷ Light gray **#97abbc**

You don't need to keep this slide in your presentation. It's only here to serve you as a design guide if you need to create new slides or download the fonts to edit the presentation in PowerPoint®




SlidesCarnival icons are editable shapes.

This means that you can:

- Resize them without losing quality.
- Change fill color and opacity.
- Change line color, width and style.

Isn't that nice? :)

Examples:



Now you can use any emoji as an icon!

And of course it resizes without losing quality and you can change the color.

How? Follow Google instructions

<https://twitter.com/googledocs/status/730087240156643328>





Free templates for all your presentation needs



For PowerPoint and
Google Slides



100% free for personal
or commercial use



Ready to use,
professional and
customizable



Blow your audience
away with attractive
visuals